

Using Experience to Improve Constrained Planning on Foliations for Multi-Modal Problems

Zachary Kingston, Constantinos Chamzas, and Lydia E. Kavraki

Abstract—Many robotic manipulation problems are *multi-modal*—they consist of a discrete set of *mode families* (e.g., whether an object is grasped or placed) each with a *continuum* of parameters (e.g., where exactly an object is grasped). Core to these problems is solving *single-mode* motion plans, i.e., given a *mode* from a mode family (e.g., a specific grasp), find a feasible motion to transition to the next desired mode. Many planners for such problems have been proposed, but complex manipulation plans may require prohibitively long computation times due to the difficulty of solving these underlying single-mode problems. It has been shown that using *experience* from similar planning queries can significantly improve the efficiency of motion planning. However, even though modes from the same family are similar, they impose different *constraints* on the planning problem, and thus experience gained in one mode cannot be directly applied to another. We present a new experience-based framework, ALEF, for such multi-modal planning problems. ALEF learns using paths from single-mode problems from a mode family, and applies this experience to novel modes from the same family. We evaluate ALEF on a variety of challenging problems and show a significant improvement in the efficiency of sampling-based planners both in isolation and within a multi-modal manipulation planner.

I. INTRODUCTION

Solving manipulation planning problems can be complex and time-consuming, as both a sequence of actions and their corresponding valid motions must be found. During search, a manipulation planner will evaluate many different variations of an action (e.g., different grasps and placements of an object), as a variation may not have a corresponding feasible motion (e.g., due to obstacles). To find valid motions, manipulation planners such as task and motion planners [1] or multi-modal planners [2], [3] use motion planning algorithms as a subroutine. Thus, improving the efficiency of motion planning improves the efficiency of manipulation planning.

Many manipulation problems are *multi-modal*, and contain a discrete set of actions (*mode families*) that are *parameterized* by continuous values, e.g., the placement of an object on a table is given by x, y -coordinates. Each parameterization of an action defines a *mode* which imposes different constraints on the problem, namely a *manifold constraint*. Moreover, as the parameters are continuous, problems with parameters “nearby” other parameters will be similar, albeit under different constraints. As each mode within a mode family defines

ZK, CC, and LEK are with the Department of Computer Science, Rice University, Houston, TX, USA [{zak, chamzas, kavraki}](mailto:{zak, chamzas, kavraki}@rice.edu) at rice.edu. ZK is supported by NASA Space Technology Research Fellowship 80NSSC17K0163. CC is supported by the NSF Graduate Research Fellowship Grant No. 1842494. This work is supported in part by NSF 1718478 and Rice University Funds.

similar problem constraints, results from prior plans could be used as *experience* to improve future queries.

Experience-based planning methods [4]–[6] learn from prior motion planning problems to expedite search in new, similar problems. For these methods, similarity is typically defined in terms of the obstacles in a changing environment, and cannot cope with changing problem constraints. In multi-modal planning, as solutions from different modes lie on disjoint manifolds, experience in one mode cannot be directly applied to another, and thus these methods cannot be used.

This work proposes a novel experience-based framework, ALEF, that can effectively reuse experiences in multi-modal problems to improve the efficiency of manipulation planning. ALEF builds a sparse roadmap within an augmented, manifold-constrained state space which unifies and relates experience gathered from different single-mode problems within a mode family. Our method learns by using paths from these problems to construct the sparse roadmap. Upon a new query, paths from the sparse roadmap are retrieved and used to bias sampling in a sampling-based planner. Learning is quick and can be run online within a manipulation planner. We demonstrate the effectiveness of our approach on challenging manipulation problems with varying environments.

II. PRELIMINARIES

In this work, we consider manipulation planning problems that are parameterized by a continuum of values. We use terminology from *multi-modal planning*, a type of manipulation planning, to describe these types of tasks (as in [2], [3]). Parameterized actions are *mode families*, where each parameterization of the action is a *mode*. We use a multi-modal planner to demonstrate our framework within the context of a manipulation planner (Sec. V-B).

Consider a robot with a configuration space \mathcal{Q} . Typically, in motion planning we are interested in finding a collision-free path σ from a point $q_{\text{start}} \in \mathcal{Q}$ to some region of interest $Q_{\text{goal}} \subset \mathcal{Q}$, where $\sigma : [0, 1] \rightarrow \mathcal{Q}_{\text{free}}$ such that $\sigma(0) = q_{\text{start}}$, $\sigma(1) \in Q_{\text{goal}}$ and $\sigma(t)$ is collision-free for all t .

A. Single-Mode Planning

A specific parameterization of an action defines a *mode* ξ which imposes a set of *manifold constraints* on a robot’s motion. Manifold constraints are determined by a *constraint function* $F^\xi : \mathbb{R}^n \rightarrow \mathbb{R}^{k_\xi}$ ($1 \leq k_\xi < n$). Constraints are satisfied when $F^\xi(q) = \mathbf{0}$; the set of all configurations which satisfy the constraint define the *mode manifold* \mathcal{M}^ξ :

$$\mathcal{M}^\xi = \{q \in \mathcal{Q} \mid F^\xi(q) = \mathbf{0}\}.$$

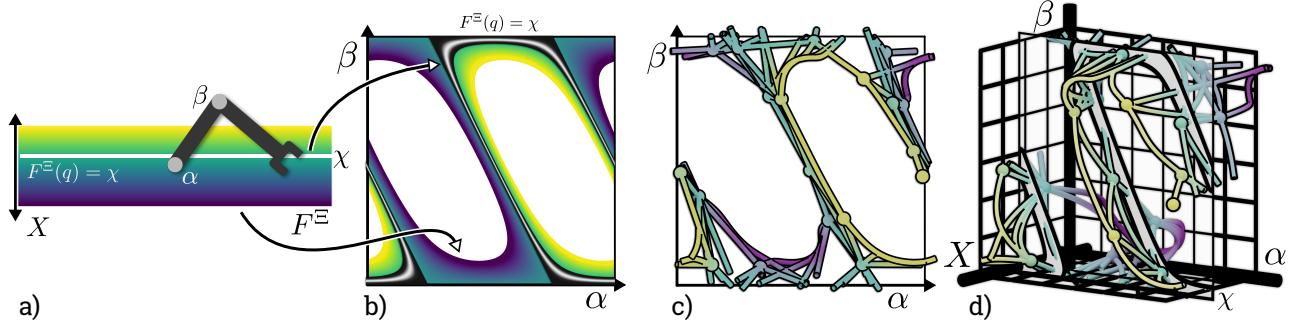


Fig. 1. Our experience-based constrained planning framework, ALEF. **a)** The arm has two continuous joints α, β . In this simple task, the arm is constrained to move its end-effector along a line, which imposes a *manifold constraint* (a mode). The line is parameterized by values from X , which determines its height (a mode family). The set of these lines forms a plane—the parameter of each line is visualized with a color gradient. An example line at $\chi \in X$ is shown in white. **b)** Visualization of the set of manifold constraints defined by this task in configuration space. The color gradient is shared with **a**—colors in the configuration space correspond to the colors in workspace (the arrows). The manifold corresponding to χ (white line) is shown in grayscale ($F^\Xi(q) = \chi$). **c)** The sparse roadmap learned by ALEF from 10 randomly sampled problems (end-effector constrained to line). The color gradient is shared with **b**, indicating each configuration’s parameters χ . Although vertices in the roadmap satisfy different constraints (modes), ALEF makes connections between them. **d)** The manifold constraint corresponding to χ intersecting our sparse roadmap. The parameters X are another dimension (see Sec. IV-A). Note that edges between vertices are geodesics on this manifold (giving their curvature), and show related experience our method can apply.

Thus, single-mode planning requires finding a valid path in \mathcal{M}^ξ , an $(n - k^\xi)$ -dimensional submanifold of \mathcal{Q} . See [7] for more on planning under manifold constraints.

B. Motion Planning within a Mode Family

We consider manipulation problems where actions have continuous parameterizations that determine how an action is done. Parameterized actions are defined as a *mode family* Ξ , as each parameterization defines a mode. For example, a robot can grasp a bar anywhere along its length: each grasp location is given by a parameter that defines a mode in the family. An example of a mode family is given in Fig. 1a.

The modes from a mode family are similar—crucially, there is continuity between the manifold constraints of a mode family’s modes. This is codified by defining mode families as *foliated manifolds* [8]. A foliated manifold is a manifold with additional structure: there exists a *transverse manifold* X^Ξ of parameters $\chi \in X^\Xi$ of dimension k , which parameterizes a set of $(n - k)$ -dimensional *leaf manifolds* (the mode manifolds) \mathcal{L}_χ for all $\chi \in X^\Xi$. A foliated manifold can also be represented as a constraint function, F^Ξ .

A mode $\xi_\chi \in \Xi$ is defined by a parameter χ , $F^\Xi(q) = \chi$. The mode’s manifold is the leaf manifold \mathcal{L}_χ :

$$\mathcal{M}^{\xi_\chi} = \mathcal{L}_\chi = \{q \in \mathcal{Q} \mid F^\Xi(q) = \chi\}.$$

An example foliated manifold and a leaf manifold is given in Fig. 1b. The definition of mode families as foliated manifolds enables us to use the general manifold-constrained motion planning framework presented in [9] used by our experience-based framework.

III. RELATED WORK

There are many kinds of manipulation planning algorithms, such as Task and Motion Planning algorithms [1], [10]–[12] and multi-modal planning algorithms [2], [3], [13]–[15]. Each algorithm handles the continuity of action parameterization differently, but all use motion planning to determine if a parameterization is valid. See [16] for a survey of techniques.

Sampling-based planners are probabilistically-complete methods able to scale to high-dimensional problems [17]–[19]. To find motions within a mode, there are many approaches for planning under manifold constraints (e.g., [9], [20]). A survey of sampling-based techniques is given in [7].

To improve efficiency, many methods adapt search *online* with information gathered during the same query. For example, the authors of [21] weight different samplers based on their performance. Other methods use prior collision checking to adapt sampling [22], improve solution quality [23], or adapt the local planner [24]. Our method too can be trained online during a single query of a manipulation planner, improving performance during search as well as on future queries.

Other experience-based methods store gathered experiences for later retrieval. Experiences can be retrieved based on start and goal similarity [4], [25] or workspace similarity [5], [6]. These methods can only retrieve experiences that satisfy the constraints they were trained on, and cannot transfer experiences between constraints. Our method uses experience from one mode and transfers that experience to other modes within the same family. The method most similar to ours is THUNDER [4], which also uses a sparse roadmap to store previous paths. However, THUNDER does not consider constraints, retrieves a path for repair rather than sampling, and requires significant processing time to store experiences.

In the context of manipulation planning, learning has been used to infer which action parameterizations are likely to be valid, both offline [26], [27] and online [3]. Other techniques use demonstrations to infer which constraints are needed to perform a task [28] or approximate constraints observed from data [29]. However, these methods learn information about the constraints themselves and do not improve motion planning. More similar to our method, the authors of [30] bias search from demonstrations to solve parameterized constrained problems. Our method learns from prior planning queries within a mode family, either in a manipulation planner or standalone, to improve performance on motion planning problems within the same mode family.

IV. THE ALEF FRAMEWORK

We present the ALEF framework (Augmented Leafs with Experience on Foliations) for experience-based manipulation planning under manifold constraints. ALEF learns using valid paths from leaf-constrained problems and applies this experience to problems constrained by different leaves within the same foliation. That is, experience from one mode can be transferred to another within the same mode family.

From valid paths (Fig. 2a), a sparse roadmap is constructed in an *augmented foliated space* (AFS) (Fig. 2b). The AFS is a manifold-constrained configuration space with respect to the foliation constraint. Configurations in the AFS are augmented to include their transverse parameters, i.e., what mode/leaf they are in. The AFS and sparse roadmap are explained in Sec. IV-A and Sec. IV-B.

Upon a new query, ALEF uses experience to bias a sampling-based algorithm. Given a start and goal configuration, nearby vertices within the sparse roadmap are retrieved. If a valid path exists in the roadmap between the retrieved start and goal vertices (Fig. 2c), the path is projected onto the current query’s leaf manifold using a projection operator (Fig. 2d). Configurations from this path that are valid are used as samples. Experience retrieval is discussed in Sec. IV-C.

A. Augmented Foliated Space (AFS)

Recall from Sec. II-B that we model mode families as foliated manifolds with constraint functions F^{Ξ} . Foliations are parameterized by a transverse manifold X , where each $\chi \in X$ corresponds to a different mode. To relate configurations across different modes, we introduce a composite space $\mathcal{Q} \times X$, where each configuration is indexed with the parameters of the mode it satisfies. This can be seen in Fig. 1c, which visualizes configurations only in \mathcal{Q} , and Fig. 1d, which visualizes the same configurations in $\mathcal{Q} \times X$.

However, not all configurations satisfy the foliation constraint. We define the *augmented foliated space* (AFS), a submanifold within the composite space $\mathcal{Q} \times X$. The AFS manifold is formally defined as:

$$\mathcal{M}^{\Xi} = \{(q, \chi) \in \mathcal{Q} \times X \mid F^{\Xi}(q) = \chi\}.$$

We leverage the general manifold-constrained sampling-based planning framework of [9] in order to model the AFS. Through the transverse dimension, connections can be made between configurations that satisfy different leaf constraints, since the AFS manifold is a superset of all leaf manifolds (Fig. 2b). The connections are such that all intermediate states satisfy the aforementioned foliation constraint.

As sampling-based methods require a metric to explore and find nearby configurations, we define a weighted metric for the AFS. This metric uses a weighted sum of the metrics of the configuration space and the transverse space. By default, we use the Euclidean metric for the transverse space. This weight relates to the relative importance of the transverse parameters versus the configuration. This can be visualized in Fig. 1d as “stretching” the X dimension. In our experiments, we weight the transverse parameter three times more than the configuration space.

B. Sparse Roadmap in the AFS

In this work, we represent experiences as valid paths gathered from leaf-constrained motion planning problems. Information from these paths is stored within a *sparse roadmap* that resides in the augmented foliated space. In particular, we employ SPARS2 [31], a method that does not require the maintenance of a dense roadmap. SPARS2 has guarantees of asymptotic near-optimality—as the roadmap “fills out” the probability of inserting a new configuration goes to zero, and paths within the roadmap are within some bound of optimal. A sparse roadmap has the benefit of finite memory requirements; the small size of the sparse roadmap, as compared to a dense roadmap, enables fast search times on new retrieval queries (Sec. IV-C). Note that we use SPARS2 within the manifold-constrained AFS. We conjecture that the theoretical properties of SPARS2 hold in this case given theoretical results from [9].

The idea of using SPARS2 as a database to store and retrieve experiences was first introduced in the THUNDER algorithm [4]. However, inserting, retrieving, and reusing paths for SPARS2 in the AFS demand different methods as compared to the standard unconstrained roadmap used in [4].

As noted by [4], a naïve insertion of the waypoints in sequential order will most likely result in the vertices of the path being disconnected within the roadmap. This stems from the way the SPARS2 chooses which vertices to connect in order to maintain sparseness. Connections between disconnected components of the roadmap are only attempted when a new vertex acts as a *connectivity* node, meaning it is “visible” from two vertices that do not belong to the same connected component. A node is visible by another if it is within a certain radius (*visibility radius*) and there is a valid connection between them. Additionally, nodes are added as *guard nodes* when there are no other nodes that are visible. However, with this insertion policy, many paths will end up disconnected. For example, when sequentially inserting a straight-line path, only nodes that act as *guards* will be added, and no connections between them will be attempted. Thus, we use the ordering heuristic proposed by [4]—first, the inserted path is interpolated at high resolution, then, nodes that are likely *guards* are inserted, and then select nodes between these guards are added as *connectivity* nodes. Additionally, to improve connectivity, we check if these newly inserted nodes can be connected to the other connected components of the roadmap. Effectively, this increases the visibility radius of SPARS2 (and thus affects SPARS2’s asymptotic near-optimality property), possibly resulting in longer paths but quickly improving roadmap connectivity. This enables ALEF to be trained online within a manipulation planner.

As the sparse roadmap is built in the AFS, edges are added between nodes from different leaves if the edge satisfies the foliation constraint and is collision-free (Fig. 2c). In the example shown in Fig. 1, nodes in the roadmap from problems that were constrained to different lines are connected by edges that correspond to a motion of the manipulator that moves between these lines. Valid connections made in the AFS are

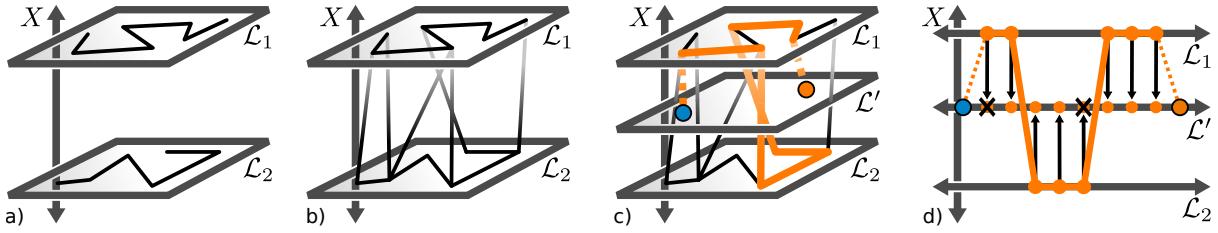


Fig. 2. Sketch of ALEF, the proposed method. **a)** Valid paths from different action parameterizations are gathered as experience. These paths exist in the leaves \mathcal{L}_1 and \mathcal{L}_2 of the foliated manifold of the action. Note that these are paths in configuration space, but are such that each configuration satisfies their leaf's constraint. **b)** ALEF builds a sparse roadmap in the augmented foliated space (AFS, Sec. IV-A), finding feasible connections across leaves. Connections are such that the foliation constraint is always satisfied. **c)** Upon a query in new leaf \mathcal{L}' , a path is found in the roadmap (highlighted) given a start and goal (blue and orange circles). **d)** The retrieved path is projected to \mathcal{L}' . Valid configurations are used as samples in the new query.

indicative of possibly valid motions on all leaves in-between two nodes, due to the continuity between the leaves of the foliated manifold. A roadmap generated by ALEF is shown in Fig. 1c. Vertices and edges are colored according to their transverse parameter. An example of this continuity can be seen in Fig. 1d, where a leaf manifold is shown intersecting the roadmap in the AFS.

C. Retrieving Experience from the AFS Roadmap

In manipulation planning, many candidate goal configurations are considered during motion planning in order to find a valid transition from one mode to the next (the intersection of manifolds). ALEF considers all goal configurations that are sampled. Given each start and goal configuration pair, their nearest neighbors in the AFS roadmap are found using the weighted AFS metric. Then, a collision-free path between these configurations within the roadmap is found using A* search (Fig. 2c). As there might be changes within the environment (e.g., changing obstacles between queries), edge validity is lazily evaluated in roadmap search, which invalidates edges as they are discovered. Paths retrieved from the roadmap hopefully contain configurations important for the current planning problem.

Given a retrieved path, the waypoints of the path are *projected* to satisfy the query's leaf constraint and are checked for collision (Fig. 2d). For projection, we use gradient descent with respect to the leaf's constraint function, but any projection operator would suffice (see [9] for further details). It is unlikely that the entire retrieved path is successfully projected onto the new leaf, thus, we cannot use the retrieved path directly. Instead, the waypoints of the retrieved path are used as samples to bias the search of a sampling-based planner. ALEF uses *all* valid waypoints from retrieved paths given *all* start and goal pairs. Samples are used with some probability $1 > \lambda > 0$. For all experiments we use $\lambda = 0.5$.

V. EXPERIMENTS

We evaluate the performance of ALEF on a “monkey” robot tasked with climbing across a set of bars and a “handoff” problem with two manipulators. Although these problems have a 2D workspace, they contain complex robots with 9 and 8 degrees-of-freedom (DOF) respectively, and are subject to non-trivial end-effector constraints. ALEF is implemented with the Open Motion Planning Library (OMPL) [9], [32].

We use PRM for all single-mode planning problems under manifold constraints. We demonstrate ALEF’s ability to learn given only a few examples and improve planning over a foliation in Sec. V-A. Then, we show how ALEF improves the performance of manipulation planning by using it within a multi-modal planner (Sec. V-B).

A. Planning in a Mode Family

The robot has 9 DOF and has two end-effectors, shown in Fig. 3a. There are two foliations we consider in this problem: all grasps of the right limb on bar 1 (the source foliation), and all grasps of the left limb on bar 2 (the destination foliation). The transverse parameter is the location on the bar the robot has grasped, shown in Fig. 3a. Problems were generated by randomly sampling grasps on bar 1, which determine the leaf of the problem. The goal is to reach the destination foliation. All problems in both the test and training sets are solvable.

Fig. 3b shows timing results for planning on 500 randomly sampled problems with a timeout of 30 seconds. Our framework achieves notable speed-up even with a few examples and continues to improve performance as the training set size increases. Additionally, even with few examples, the variance in solution time decreases, showing that our framework learns to solve “hard” problems faster. This is also visible in Fig. 3c, which shows the cumulative distribution of solving the planning problem versus planning time.

Fig. 3d shows the path retrieval and valid state ratio distributions for our framework over the 500 tested queries. The path retrieval ratio is the ratio of how many paths were successfully retrieved for all start/goal query pairs. A ratio of 1 means that all queries had relevant experience retrieved from the roadmap, while 0 means that no relevant experience was found. The valid state ratio is the ratio of the states from retrieved paths that were successfully projected onto to the new leaf. A high valid state ratio means that the experience retrieved was “useful” to the new problem. Even with only 10 plans inserted into the roadmap, a high ratio of experience is retrieved. However, the average of the ratio of valid states was low (the peak at 0 in Fig. 3d). As the amount of experience in the roadmap increases so does the ratio of the valid states, improving the performance of ALEF.

B. Multi-Modal Planning

Our framework also improves the performance of a manipulation planner, where there are multiple mode families

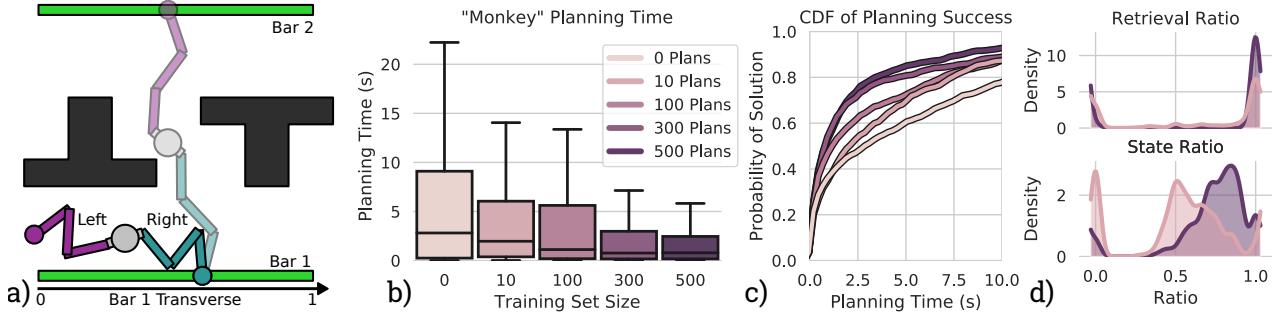


Fig. 3. a) The “monkey” environment. A two-limbed robot with 9 DOF must grasp Bar 2. Start configurations are sampled such that the robot grasps Bar 1. The location of the robot’s grasp on the bar corresponds to the transverse parameter of the foliation, indicated below the bar. b) Timing results for 500 trials. On the x -axis, our method is trained on increasing sizes of training plans. Significant speed-up is achieved given only a few examples. c) Cumulative probability of finding a solution versus time. Our method solves more problems faster as experience is gathered. d) Retrieval ratio and valid state ratio for ALEF trained on 10 and 500 examples. As ALEF is trained, more paths are retrieved and the amount of relevant experience increases.

to traverse in a single query. ALEF was implemented in the multi-modal planner of [3]. We test the following variations:

- “None”: This is the baseline motion planner in [3] and does not utilize our framework at all.
- “Adaptive”: Here, ALEF learns *online* while the multi-modal planner is running. The results of every planning query the multi-modal planner makes are inserted into the roadmap. ALEF can only learn during the current multi-modal planning query. Learning time is included.
- “ n Plans”: Here, ALEF was trained using all motion plans generated from n multi-modal planning queries.

1) *Monkey Example*: Fig. 4a shows an extension of the environment from Fig. 3a with three bars (6 foliations). The robot is tasked with climbing to reach a goal past the far bar. That is, the planner must find a sequence of feasible mode transitions to reach the goal, requiring at least three mode transitions (e.g., from grasping the initial bar to grasping the middle, then middle to far, and then far to the goal). The monkey starts always from the same configuration, but obstacles are varied between queries. Specifically, each obstacle can rotate 20 degrees about its center and vary in position by as much as its thinnest width. Depending on the pose of the obstacles, an alternate route to the goal might need to be taken, as the middle corridor might be closed.

Timing results for total multi-modal planning time are presented in Fig. 4b and Fig. 4c. Starting from nothing, the “Adaptive” planner provides a small benefit over baseline performance, showing that our framework helps solve queries even with limited experience. Note that these reported times include training time for ALEF, which is negligible. Offline training from other queries gives substantial benefit and accelerates multi-modal planning.

Fig. 4d shows path retrieval and valid state ratio distributions for all single-mode plans made by the multi-modal planner. Here, “Adaptive” does not retrieve much experience, and the experience it retrieves is typically unhelpful, given the low valid state ratio. As our method is trained on more plans, the retrieval ratio and valid state ratio both increase significantly, which is corroborated by their performance.

2) *Handoff Example*: Fig. 5a shows a “handoff” environment (8 DOF), where an object must be transferred from one

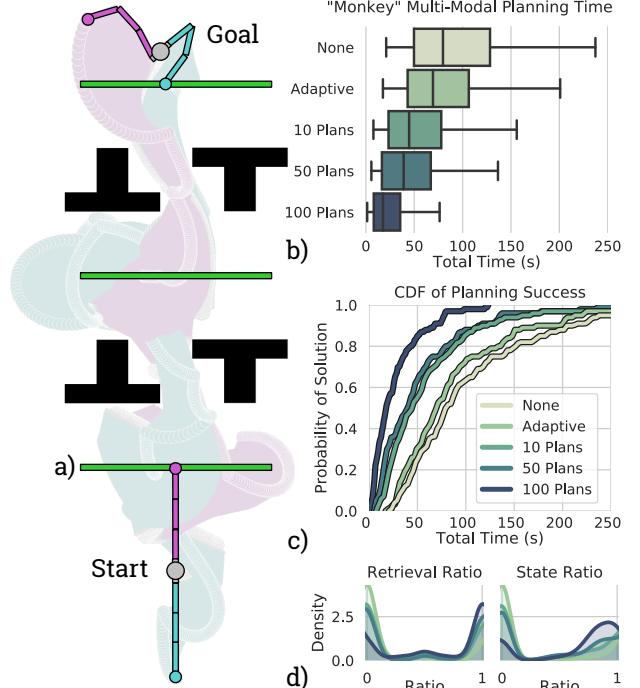


Fig. 4. a) The “monkey” environment. The start and an example goal configuration are highlighted. The robot must alternate grasps to climb from bar to bar to reach the goal. The swept volume of an example multi-modal plan is shown. b) Total multi-modal planning time for 100 trials for each method, including training time for ALEF. Obstacles vary in position and rotation between each query. Shown are a baseline method with no experience (“None”), ALEF which learns from scratch (“Adaptive”), and ALEF trained on n multi-modal queries (“ n plans”). Starting from scratch, our method provides a small benefit over the baseline. Training from prior queries provides substantial benefit. c) Timing results presented as a cumulative distribution curve. d) Retrieval and valid state ratios for all motion planning queries. ALEF retrieves more relevant experience as its training set increases.

side to the other (stills 1 and 3 in Fig. 5a). However, due to the length of the manipulators and the obstacle in the middle, the object must be handed off (still 2 in Fig. 5a). Additionally, the end-effector of the manipulator is constrained to always remain upright. Here, there is a mode family for grasping the object anywhere along its length, and another mode family for placing the object anywhere on the flat surface. Similar to before, the problem starts from the same configuration, but the gray obstacle varies between queries. Here, the gray

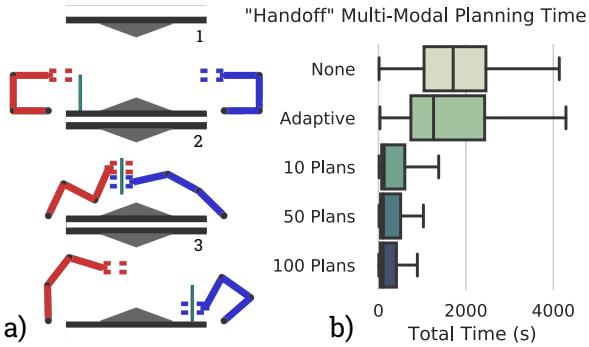


Fig. 5. a) Three stills from a multi-modal plan in the “handoff” environment. The thin object must be handed off so it can be transferred to the other side of the environment. Many motion plans are attempted in this environment, as the initial grasp mode determines if a handoff is possible. b) Multi-modal planning time over 100 trials for each method, including training time for ALEF. Obstacles vary in position and rotation between each multi-modal query. ALEF shows significant improvement over baseline.

obstacle can vary up to ± 20 degrees about its center, and in position by the height of its peak in both the x - and y -axes.

Results for total multi-modal planning time are presented in Fig. 4b. As before, the “Adaptive” planner provides a small benefit over baseline, while training ALEF with multiple queries gives dramatic performance improvements. This example reveals the generality of our framework.

VI. CONCLUSION

We have presented our framework, ALEF, for experience-based planning in the context of manipulation planning. ALEF transfers experience between a continuum of manifold-constrained problems, specifically problems that are drawn from a “mode family”, or foliation. ALEF builds a sparse roadmap in an augmented space that makes connections between problems with different constraints and uses this roadmap to retrieve experience as a sampler on future queries. ALEF provides significant speedup in isolation and within a manipulation planner given only a few examples. In the future, we plan to demonstrate ALEF on problems with 3D workspaces using realistic robots, as well as foliations with higher dimensional traverse manifolds. Moreover, we plan to investigate other experience storage methods and apply our method within other manipulation planning algorithms.

REFERENCES

- [1] N. T. Dantam, Z. Kingston, S. Chaudhuri, and L. E. Kavraki, “An incremental constraint-based framework for task and motion planning,” *Int. J. of Robotics Research*, vol. 37, no. 10, pp. 1134–1151, 2018.
- [2] K. Hauser and V. Ng-Thow-Hing, “Randomized multi-modal motion planning for a humanoid robot manipulation task,” *Int. J. of Robotics Research*, vol. 30, no. 6, pp. 678–698, 2011.
- [3] Z. Kingston, A. M. Wells, M. Moll, and L. E. Kavraki, “Informing multi-modal planning with synergistic discrete leads,” in *IEEE Int. Conf. Robot. Autom.*, 2020, to appear.
- [4] D. Coleman, I. A. Sucan, M. Moll, K. Okada, and N. Correll, “Experience-based planning with sparse roadmap spanners,” in *IEEE Int. Conf. Robot. Autom.*, 2015, pp. 900–905.
- [5] C. Chamzas, A. Shrivastava, and L. E. Kavraki, “Using local experiences for global motion planning,” in *IEEE Int. Conf. Robot. Autom.*, May 2019, pp. 8606–8612.
- [6] N. Jetchev and M. Toussaint, “Fast motion planning from experience: trajectory prediction for speeding up movement generation,” *IEEE J. Robot. Autom.*, vol. 34, no. 2, pp. 111–127, 2013.
- [7] Z. Kingston, M. Moll, and L. E. Kavraki, “Sampling-based methods for motion planning with constraints,” *Annual Review of Control, Robot., and Autom. Syst.*, vol. 1, no. 1, pp. 159–185, 2018.
- [8] M. Spivak, *A Comprehensive Introduction to Differential Geometry*. Publish or Perish, 1999.
- [9] Z. Kingston, M. Moll, and L. E. Kavraki, “Exploring implicit spaces for constrained sampling-based planning,” *Int. J. of Robotics Research*, vol. 38, no. 10–11, pp. 1151–1178, 2019.
- [10] L. P. Kaelbling and T. Lozano-Pérez, “Hierarchical task and motion planning in the now,” in *IEEE Int. Conf. Robot. Autom.*, 2011, pp. 1470–1477.
- [11] S. Srivastava, E. Fang, L. Riano, R. Chitnis, S. Russell, and P. Abbeel, “Combined task and motion planning through an extensible planner-independent interface layer,” in *IEEE Int. Conf. Robot. Autom.*, 2014, pp. 639–646.
- [12] C. R. Garrett, T. Lozano-Pérez, and L. P. Kaelbling, “Sampling-based methods for factored task and motion planning,” *Int. J. of Robotics Research*, vol. 37, no. 13–14, pp. 1796–1825, 2018.
- [13] J. Barry, L. P. Kaelbling, and T. Lozano-Pérez, “A hierarchical approach to manipulation with diverse actions,” in *IEEE Int. Conf. Robot. Autom.*, 2013, pp. 1799–1806.
- [14] W. Vega-Brown and N. Roy, “Asymptotically optimal planning under piecewise-analytic constraints,” in *Int. Wksp. on the Algorithmic Foundations of Robotics*. Springer, 2016.
- [15] P. S. Schmitt, F. Wirnhofer, K. M. Wurm, G. von Wichert, and W. Burgard, “Modeling and planning manipulation in dynamic environments,” in *IEEE Int. Conf. Robot. Autom.*, 2019, pp. 176–182.
- [16] C. R. Garrett, R. Chitnis, R. Holladay, B. Kim, T. Silver, L. P. Kaelbling, and T. Lozano-Pérez, “Integrated task and motion planning,” *Annual Review of Control, Robot., and Autom. Syst.*, vol. 4, no. 1, 2021.
- [17] H. M. Choset, S. Hutchinson, K. M. Lynch, G. Kantor, W. Burgard, L. E. Kavraki, and S. Thrun, *Principles of Robot Motion: Theory, Algorithms, and Implementation*. MIT Press, 2005.
- [18] S. M. LaValle, *Planning Algorithms*. Cambridge University Press, 2006.
- [19] L. E. Kavraki and S. M. LaValle, *Springer Handbook of Robotics*, 2nd ed. Springer, 2016, ch. Motion Planning, pp. 139–162.
- [20] D. Berenson, S. S. Srinivasa, and J. J. Kuffner, “Task space regions: A framework for pose-constrained manipulation planning,” *Int. J. of Robotics Research*, vol. 30, no. 12, pp. 1435–1460, Oct. 2011.
- [21] D. Hsu, G. Sánchez-Ante, and Z. Sun, “Hybrid PRM sampling with a cost-sensitive adaptive strategy,” in *IEEE Int. Conf. Robot. Autom.*, 2005, pp. 3874–3880.
- [22] B. Burns and O. Brock, “Toward optimal configuration space sampling,” in *Robotics: Science and Syst.*, 2005, pp. 105–112.
- [23] S. S. Joshi and T. Panagiotis, “Non-parametric informed exploration for sampling-based motion planning,” in *IEEE Int. Conf. Robot. Autom.*, 2019, pp. 5915–5921.
- [24] T. Lai, P. Moreira, F. Ramos, and G. Francis, “Bayesian local sampling-based planning,” vol. 5, no. 2, pp. 1954–1961, April 2020.
- [25] D. Berenson, P. Abbeel, and K. Goldberg, “A robot path planning framework that learns from experience,” in *IEEE Int. Conf. Robot. Autom.*, 2012, pp. 3671–3678.
- [26] Z. Wang, C. R. Garrett, L. P. Kaelbling, and T. Lozano-Pérez, “Active model learning and diverse action sampling for task and motion planning,” in *IEEE/RSJ Int. Conf. on Intell. Robots and Syst.*, 2018, pp. 4107–4114.
- [27] B. Kim, Z. Wang, L. P. Kaelbling, and T. Lozano-Pérez, “Learning to guide task and motion planning using score-space representation,” *Int. J. of Robotics Research*, vol. 38, no. 7, pp. 793–812, 2019.
- [28] G. Ye and R. Alterovitz, “Guided motion planning,” in *Springer Tracts in Advanced Robot.* Springer, 2017, pp. 291–307.
- [29] G. Sutanto, I. M. R. Fernández, P. Englert, R. K. Ramachandran, and G. S. Sukhatme, “Learning equality constraints for motion planning on manifolds,” in *Conf. on Robot Learning*, 2020.
- [30] M. Phillips, V. Hwang, S. Chitta, and M. Likhachev, “Learning to plan for constrained manipulation from demonstrations,” *Autonomous Robots*, vol. 40, no. 1, pp. 109–124, 2016.
- [31] A. Dobson and K. E. Bekris, “Improving sparse roadmap spanners,” in *IEEE Int. Conf. Robot. Autom.*, 2013, pp. 4106–4111.
- [32] I. A. Sucan, M. Moll, and L. E. Kavraki, “The Open Motion Planning Library,” *IEEE Robot. Autom. Magazine*, vol. 19, no. 4, pp. 72–82, 2012.